



AI RELIGIOUS KNOWLEDGE EMPLOYEE

FaithBrains

Cited, scholar-reviewed consultation of Quran and Hadith for learning. Consultation, not religious authority.

Document 11 · FaithBrains

Audience Product, engineering, GTM & leadership

Status Plan · **Version** 1.0

Classification Confidential — Internal

01

02

03

04

05

06

07

08

09

10

11

12

13

14

15

16

17

18

19



01Executive Summary

FaithBrains is an AI religious-knowledge employee on AvatarOS for learning and consultation of Islamic sources, built around strict governance: every answer is sourced and cited, scholars review content, and the system never issues rulings (fatwa) or presents itself as religious authority. It supports Quran consultation, Hadith consultation, and general Islamic learning, deferring to qualified human scholars for matters of ruling.

3 modules	AvatarOS fully reused core	MVP→LT 4-phase roadmap	SaaS subscription model
---------------------------------	--	--	---

Positioning — AI Religious Knowledge Employee

Cited, scholar-reviewed consultation of Quran and Hadith for learning. Consultation, not religious authority.

02 Vision, Mission & Modules

Vision. Make FaithBrains the default ai religious knowledge employee for any organization — capable, trustworthy, and instantly deployable on AvatarOS.

Mission. Deliver FaithBrains's modules as a coherent ai religious knowledge employee customers can trust.

Modules:

Module	What it does
Quran Consultation	Cited consultation of Quranic verses with translation provenance and tafsir references; presents scholarly views, not rulings.
Hadith Consultation	Hadith lookup with collection, grading, and chain references where available; flags authenticity context.
Islamic Learning Assistant	Structured, plain-language learning with citations and explicit deferral to scholars on rulings.

03 Business Opportunity

FaithBrains addresses concrete, recurring pain with measurable ROI. The opportunity is amplified because AvatarOS removes most build cost.

Pain / opportunity	How FaithBrains solves it
Hard to find sourced answers	Cited consultation with provenance
Misinformation online	Scholar-reviewed, transparent sourcing
Language/access barriers	Multilingual learning assistant
Learners want structured study	Guided learning paths

04 User Personas

Persona	Goal	What they value
Learner	Understand sources accurately	Cited, plain-language consultation
Student of knowledge	Study with references	Source-linked Hadith/Quran lookup
Educator/Imam	Teaching aid	Scholar-reviewed content, citations
New Muslim	Gentle, accurate intro	Learning assistant + safeguards

05

Market & Competitor Analysis

Competitors tend to own one slice; FaithBrains wins by combining capabilities on the shared AvatarOS core.

Competitor	They offer	FaithBrains advantage
Quran apps	Text + search	Conversational, cited, scholar-reviewed
General chatbots	Unsourced answers	Strict citation + scholar review + no fatwa
Islamic Q&A; sites	Human-curated	Real-time, multilingual, transparent sourcing

06

Product Strategy & Platform Leverage

FaithBrains is built entirely on AvatarOS (document 02). It reuses platform capabilities and adds only a thin product layer — this is what makes it fast to build and cheap to run.

How FaithBrains leverages AvatarOS (no architecture duplication)
Reused from platform: Knowledge consultation + citation engine (central here), RAG, freshness/authority model, voice + avatar, shared memory (learner path), tenancy.
Net-new in this product: Vetted source corpus (translations, tafsir, hadith collections with grading), scholar-review workflow, fatwa/ruling refusal guard, multilingual presentation.

AvatarOS layers consumed:

AvatarOS Layer	Responsibility & provider posture (3rd-party → indigenous)
L1 — Ingress	Speech recognition (Deepgram→indigenous), VAD, barge-in, telephony/WebRTC/SIP
L2 — Cognition	LLM router (OpenAI/Anthropic→indigenous), prompt assembly, tool-calling, guardrails
L3 — Memory	Working, episodic, semantic & long-term memory; vector + graph store; consolidation
L4 — Knowledge	RAG, knowledge consultation, citation engine, freshness & provenance
L5 — Orchestration	Agent runtime, multi-agent collaboration, planner/executor, skill registry
L6 — Affect	Emotional intelligence (Hume→indigenous), personality system, prosody control
L7 — Egress	Voice synthesis (ElevenLabs/Cartesia→indigenous), avatar rendering (HeyGen→indigenous)
L8 — Realtime	Streaming pipeline, sub-700ms latency budget, turn-taking, interruption handling
L9 — Platform	Multi-tenant control plane, billing/metering, identity, observability, eval harness

07

Technical & Multi-Tenant Architecture

FaithBrains runs as a tenant-scoped product on AvatarOS. It inherits multi-tenant isolation, billing/metering, identity, and the realtime pipeline; it contributes a product service holding its domain logic.

```
customer tenant --> FaithBrains product service (thin)
> AvatarOS realtime: ASR > LLM > memory > knowledge > agents > affect > TTS+avatar
> control plane: tenant isolation . billing . identity . observability (inherited)
> net-new (this product only):
  Vetted source corpus (translations, tafsir, hadith collections with grading),
  scholar-review workflow, fatwa/ruling refusal guard, multilingual presentation.
```

No architecture duplication

FaithBrains does not re-implement memory, voice, avatar, orchestration, tenancy, or billing. Those are AvatarOS shared services. Only domain-specific logic is new.

08

AI Agent Design

FaithBrains uses the AvatarOS planner/agents/critic pattern, configured for its domain. The critic encodes this product's most important guardrail.

```
PLANNER: intent (verse|hadith|learn) – never 'ruling'  
AGENTS : source-retriever(cited) · translator/explainer · scholar-review queue  
CRITIC : AUTHORITY GUARD – no fatwa/rulings; cite sources; defer to scholars; flag sensitive topics  
MEMORY : learner path (episodic) ► topics studied (semantic)
```

09 Memory Design

FaithBrains configures the AvatarOS shared memory layer for its domain — it chooses what is salient, not how memory works.

Memory type	What FaithBrains stores
Working	Current question + retrieved sources
Episodic	Learner's study sessions
Semantic	Topics covered, learning path
Long-term	Learner goals & level (no profiling of belief)

10

User Journey & Workflows

User journey (textual diagram):

```
learner asks ▶ retrieve cited sources (Quran/Hadith/tafsir)
▶ present with provenance + scholarly views ▶ authority critic ensures no ruling
▶ defer to scholar where needed ▶ scholar-reviewed content surfaced
```

Core workflow (textual diagram):

```
scholars vet corpus ▶ content marked reviewed ▶ learner asks
▶ cited retrieval ▶ critic blocks ruling-style output, adds disclaimers
▶ sensitive/edge topics ▶ scholar-review queue ▶ transparent sourcing shown
```

11

UX Direction

FaithBrains uses the shared CTB design system with a persona tuned to its domain. Priorities:

- Latency-honest states (listening / thinking / speaking) with text fallback.
- Trust affordances: visible citations and correction where the domain demands accuracy.
- Frictionless escalation or hand-off to a human when confidence is low.
- Accessibility: captions, adjustable pace, keyboard navigation.

12

Database & API Concepts

Product entities (extend AvatarOS core entities from doc 02 — Tenant, Persona, MemoryRecord, UsageEvent — never fork them):

Entity	Key fields	Notes / relationships
Source	id,type(quran hadith tafsir),ref,translation,grading	Provenance + authority
ScholarReview	id,content_id,scholar,status,notes	Governance
Learner	id,path,topics	Study progress
EscalationFlag	id,query,reason	Routes to human scholars

API surface (product endpoints over the AvatarOS internal contracts):

- POST /faithbrains/session — start a domain session (tenant-scoped).
- WS /stream — realtime audio/text (inherited from platform).
- GET /faithbrains/records — domain records (e.g., Source).
- GET /usage — metering for billing (inherited).

13

Security, Compliance & Deployment

FaithBrains inherits AvatarOS security: encryption, RBAC, tenant isolation, PII tagging, audit, and selectable compliance modes (doc 02). Deployment is cloud-native and region-aware with vendor fallback.

- Tenant-scoped data; no cross-tenant access; full audit trail.
- Compliance modes selectable per tenant (e.g., residency; HIPAA-capable for regulated buyers).
- Every external provider has a degrade path; no hard fail on vendor outage.

Governance & safeguards (binding)

All content is sourced and cited; scholars review content before it is surfaced; the system never issues rulings and defers to qualified human scholars; sensitive topics route to human review; sourcing and provenance are always transparent to the user.

14

Scaling & Monetization

Scales on AvatarOS autoscaling (concurrent sessions, minutes, tokens). Monetized as a SaaS subscription with metered overage.

Plan	Price	Includes	Target
Free	\$0	Cited consultation, capped	Individuals
Supporter	\$5/mo	Unlimited, learning paths	Engaged learners
Institution	Custom	Mosques/schools, admin, content	Organizations
Enterprise	Custom	Custom corpus, scholar network, SLAs	Large institutions

Monetization recommendations:

- Land on a low-friction tier; expand by seats and modules.
- Meter heavy resources (voice minutes, avatar seconds) to protect margin.
- Use AvatarOS indigenous-tech roadmap (doc 01) to improve gross margin over time.

15

Roadmap, Milestones & Launch Sequence

Phased delivery, leaning on the platform so the MVP is reachable quickly.

Phase	Timeline	Scope
MVP	0–3 mo	Core flagship module live on AvatarOS (Quran Consultation), single tenant-ready
V1	3–6 mo	Add Hadith Consultation; integrations; analytics; multi-tenant GA
V2	6–12 mo	Remaining modules (Islamic Learning Assistant); enterprise readiness
Long-term	12 mo+	Indigenous-tech cost gains; deeper autonomy within guardrails; portfolio cross-sell

Launch sequence:

- 1 Private beta with 3–5 design partners on the flagship module.
- 2 Instrument KPIs; harden the product-specific guardrail (critic).
- 3 Public launch of MVP; reference story for AvatarOS API (doc 03).
- 4 Expand modules; pursue enterprise compliance; cross-sell within CTB portfolio.

16 Team Structure

FaithBrains runs as a small squad because AvatarOS does the heavy lifting. Platform, research, and trust teams (docs 01–02) are shared, not duplicated here.

Role	FTE	Focus
Product Lead	1	Roadmap, GTM alignment, persona ownership
Conversation Designer	1	Dialogue, persona, prompt + eval design
Backend / Platform Eng	2	AvatarOS integration, tenant config, APIs
ML / Applied Scientist	1	Retrieval, memory tuning, eval benchmarks
Frontend / Avatar Eng	1	Embed widget, avatar UX, accessibility
QA / Eval	0.5	Conversation QA, regression eval harness
Scholar Review Board	—	Content vetting & governance (advisory)

17 KPIs, Success Metrics & Risks

KPIs & success metrics:

KPI	Definition	Target
Citation coverage	Answers with valid sources	100%
Ruling-refusal adherence	No fatwa issued	100%
Scholar-review coverage	Content reviewed	>95%
Sensitive-topic escalation	Correctly routed	100%

Key risks & mitigations:

Risk	Likelihood	Impact	Mitigation
Issues a religious ruling (fatwa)	Medium	High	Hard authority guard; refuse + defer to scholars; disclaimers
Misattributed/weak Hadith	Medium	High	Grading metadata; scholar review; show authenticity context
Sectarian bias / offense	Medium	High	Present scholarly views neutrally; scholar oversight; transparency
Translation/tafsir inaccuracy	Medium	High	Vetted sources only; provenance shown; review workflow

18

Detailed Task Breakdown

Indicative engineering tasks to reach MVP, organized by area. Platform tasks are configuration, not construction.

Area	Tasks (MVP)
Platform config	Provision tenant, select persona/voice/avatar, wire memory scope, enable metering
Product service	Domain models (Source, ScholarReview, Learner, EscalationFlag), state machine, Quran Consultation logic
Guardrail (critic)	Implement & red-team the product-specific critic (issues a religious ruling (fatwa))
Integrations	Vetted source corpus (translations, tafsir, hadith collections with grading), scholar-review workflow, fatwa/ruling refusal guard, multilingual presentation.
UX	Embed widget, transcript/citation UI, escalation/hand-off flow
Eval	Domain eval set, KPI instrumentation, regression harness

19

Future Evolution

From consultation to a transparent, scholar-governed learning companion across multiple faith corpora — always positioned as a sourced learning aid that defers to qualified human scholars, never as religious authority.

Platform compounding

As AvatarOS research (doc 01) lands indigenous latency/memory/affect, FaithBrains gets faster, cheaper, and more capable with no rewrite — the core advantage of building every product on one platform.